

```
In [0]: 1 from google.colab import drive
        2 drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Type *Markdown* and LaTeX: α^2

Colab uses GPU

```
In [3]: 1 #' ' means CPU whereas '/device:G:0' means GPU
        2 import tensorflow as tf
        3 tf.test.gpu_device_name()
```

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade \(https://www.tensorflow.org/guide/migrate\)](https://www.tensorflow.org/guide/migrate) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow_version 1.x` magic: [more info \(https://colab.research.google.com/notebooks/tensorflow_version.ipynb\)](https://colab.research.google.com/notebooks/tensorflow_version.ipynb).

```
Out[3]: '/device:GPU:0'
```

Testing GPU utilization

```
1
```

```
In [4]: 1 # memory footprint support libraries/code
2 !ln -sf /opt/bin/nvidia-smi /usr/bin/nvidia-smi
3 !pip install gputil
4 !pip install psutil
5 !pip install humanize
6 import psutil
7 import humanize
8 import os
9 import GPUutil as GPU
10 GPUs = GPU.getGPUs()
11 print(GPUs)
12 #XXX: only one GPU on Colab and isn't guaranteed
13 gpu = GPUs[0]
14 def printm():
15     process = psutil.Process(os.getpid())
16     print("Gen RAM Free: " + humanize.naturalsize( psutil.virtual_memory()))
17     print("GPU RAM Free: {0:.0f}MB | Used: {1:.0f}MB | Util {2:3.0f}% | Total {3:.0f}MB".format(
18         psutil.globals['nvidia'].nvidia_smi.commands['smi'].get_free_memory(gpu),
        psutil.globals['nvidia'].nvidia_smi.commands['smi'].get_used_memory(gpu),
        psutil.globals['nvidia'].nvidia_smi.commands['smi'].get_utilization(gpu),
        psutil.globals['nvidia'].nvidia_smi.commands['smi'].get_total_memory(gpu)))
18 printm()
```

Collecting gputil

Downloading <https://files.pythonhosted.org/packages/ed/0e/5c61eedde9f6c87713e89d794f01e378cfd9565847d4576fa627d758c554/GPUutil-1.4.0.tar.gz> (<http://files.pythonhosted.org/packages/ed/0e/5c61eedde9f6c87713e89d794f01e378cfd9565847d4576fa627d758c554/GPUutil-1.4.0.tar.gz>)

Building wheels for collected packages: gputil

Building wheel for gputil (setup.py) ... done

Created wheel for gputil: filename=GPUutil-1.4.0-cp36-none-any.whl size=7413 sha256=c177abcedcb0977a7b4efb71616bf761a3bad716abb85b2327a42fcc6894a064

Stored in directory: /root/.cache/pip/wheels/3d/77/07/80562de4bb0786e5ea186911a2c831fdd0018bda69beab71fd

Successfully built gputil

Installing collected packages: gputil

Successfully installed gputil-1.4.0

Requirement already satisfied: psutil in /usr/local/lib/python3.6/dist-packages (5.4.8)

Requirement already satisfied: humanize in /usr/local/lib/python3.6/dist-packages (0.5.1)

[<GPUutil.GPUutil.GPU object at 0x7f94c0723358>]

Gen RAM Free: 26.2 GB | Proc size: 457.5 MB

GPU RAM Free: 11372MB | Used: 69MB | Util 1% | Total 11441MB

Import Necessary Libraries

```
In [57]: 1 import os
          2 import zipfile
          3 import numpy as np
          4 import pandas as pd
          5 import seaborn as sns
          6 import tensorflow as tf
          7 from tensorflow import keras
          8 import matplotlib
          9 import matplotlib.pyplot as plt
         10
         11 from keras_preprocessing.image import ImageDataGenerator, load_img
         12 from keras.models import Sequential
         13 from keras.layers import Dense, Dropout, Flatten, Activation, BatchNorm
         14 from keras.layers.convolutional import Conv2D
         15 from keras.layers.convolutional import MaxPooling2D
         16 print(tf.__version__)
         17
```

1.15.0

Unzip the Dataset

```
In [0]: 1 zip_ref = zipfile.ZipFile('drive/My Drive/Datasets/Flowers.zip', 'r')
          2 zip_ref.extractall('Flower') # unzip directory
          3 zip_ref.close()
```

Load Data and Create Data Frame

```
In [5]: 1 flower_class = pd.read_csv('Flower/classlabels.txt', header = None, na
        2 flower_class
```

```
Out[5]:
```

	Images	Class
0	JFT_00001.jpg	1
1	JFT_00002.jpg	1
2	JFT_00003.jpg	1
3	JFT_00004.jpg	1
4	JFT_00006.jpg	1
...
1474	JFT_01467.jpg	29
1475	JFT_01470.jpg	30
1476	JFT_01472.jpg	30
1477	JFT_01476.jpg	30
1478	JFT_01477.jpg	30

1479 rows × 2 columns

```
In [6]: 1 flower_class['N_Images'] = flower_class['Images'].str.strip('JFT_0000').
        2 flower_class
```

```
Out[6]:
```

	Images	Class	N_Images
0	JFT_00001.jpg	1	1
1	JFT_00002.jpg	1	2
2	JFT_00003.jpg	1	3
3	JFT_00004.jpg	1	4
4	JFT_00006.jpg	1	6
...
1474	JFT_01467.jpg	29	1467
1475	JFT_01470.jpg	30	1470
1476	JFT_01472.jpg	30	1472
1477	JFT_01476.jpg	30	1476
1478	JFT_01477.jpg	30	1477

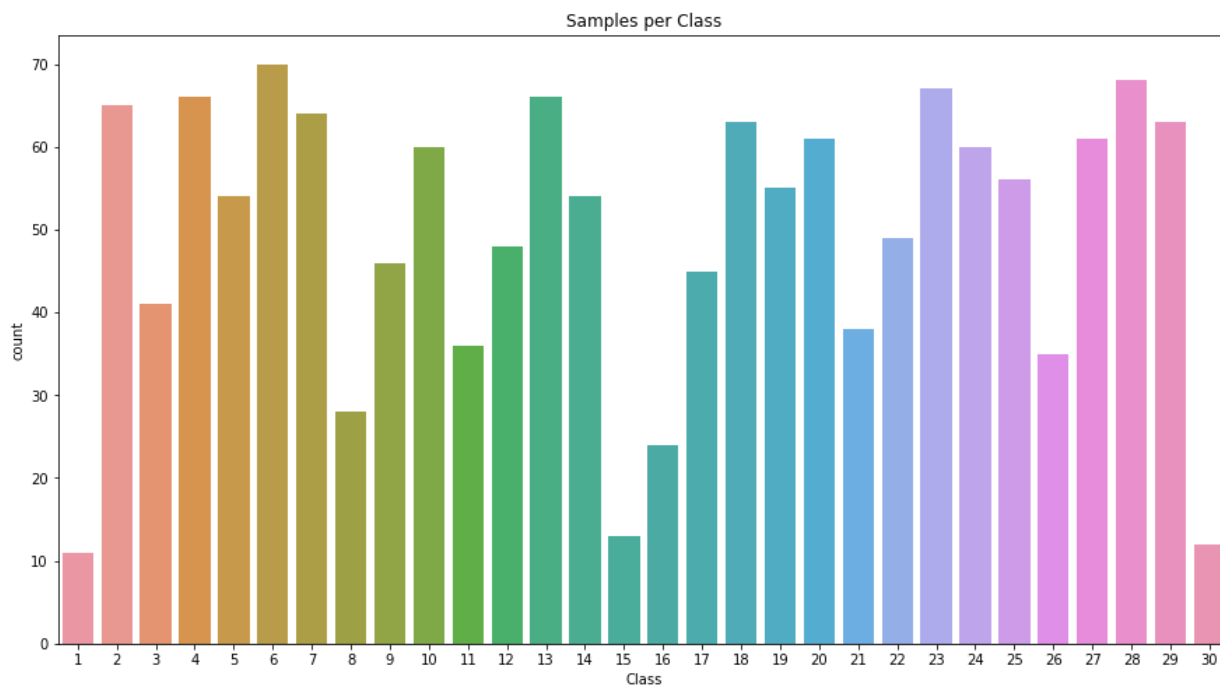
1479 rows × 3 columns

Calculate the Baseline

```
In [7]: 1 base_line = 100*flower_class['Class'].value_counts()/len(flower_class)
        2 base_line
```

```
Out[7]: 6      4.732928
        28      4.597701
        23      4.530088
        13      4.462475
         4      4.462475
         2      4.394861
         7      4.327248
        18      4.259635
        29      4.259635
        20      4.124408
        27      4.124408
        24      4.056795
        10      4.056795
        25      3.786342
        19      3.718729
        14      3.651116
         5      3.651116
        22      3.313049
        12      3.245436
         9      3.110210
        17      3.042596
         3      2.772143
        21      2.569304
        11      2.434077
        26      2.366464
         8      1.893171
        16      1.622718
        15      0.878972
        30      0.811359
         1      0.743746
        Name: Class, dtype: float64
```

```
In [12]: 1 plt.figure(figsize=(15, 8))
2         sns.countplot(x='Class',data=flower_class)
3         plt.title("Samples per Class")
4         plt.show()
```



```
In [13]: 1 train1 = pd.read_csv('Flower/tst1.txt',header= None,names=[ 'N_Images' ])
2         train1['N_Images'] = train1['N_Images'].apply(str)
3         train1df = flower_class.merge(train1, on= 'N_Images')
4         train1df['Class'] = train1df['Class'].apply(str)
5         train1df
6
7
```

Out[13]:

	Images	Class	N_Images
0	JFT_00005.jpg	1	5
1	JFT_00007.jpg	1	7
2	JFT_00011.jpg	1	11
3	JFT_00012.jpg	2	12
4	JFT_00013.jpg	2	13
...
916	JFT_01467.jpg	29	1467
917	JFT_01470.jpg	30	1470
918	JFT_01472.jpg	30	1472
919	JFT_01476.jpg	30	1476
920	JFT_01477.jpg	30	1477

921 rows × 3 columns

```
In [14]: 1 val1 = pd.read_csv('Flower/val1.txt', sep = ' ', header= None, names= ['N_I
2 val1['N_Images'] = val1['N_Images'].apply(str)
3 valldf = flower_class.merge(val1, on= 'N_Images')
4 valldf['Class'] = valldf['Class'].apply(str)
5 valldf
```

Out[14]:

	Images	Class	N Images
0	JFT_00001.jpg	1	1
1	JFT_00003.jpg	1	3
2	JFT_00006.jpg	1	6
3	JFT_00008.jpg	1	8
4	JFT_00018.jpg	2	18
...
274	JFT_01462.jpg	29	1462
275	JFT_01469.jpg	30	1469
276	JFT_01471.jpg	30	1471
277	JFT_01474.jpg	30	1474
278	JFT_01479.jpg	30	1479

279 rows × 3 columns

```
In [15]: 1 test1 = pd.read_csv('Flower/trn1.txt', sep = ' ', header= None, names=[ 'N_
2 test1['N_Images'] = test1['N_Images'].apply(str)
3 test1df = flower_class.merge(test1, on= 'N_Images')
4 test1df['Class'] = test1df['Class'].apply(str)
5 test1df
```

```
Out[15]:
```

	Images	Class	N_Images
0	JFT_00002.jpg	1	2
1	JFT_00004.jpg	1	4
2	JFT_00009.jpg	1	9
3	JFT_00010.jpg	1	10
4	JFT_00021.jpg	2	21
...
274	JFT_01466.jpg	29	1466
275	JFT_01468.jpg	30	1468
276	JFT_01473.jpg	30	1473
277	JFT_01475.jpg	30	1475
278	JFT_01478.jpg	30	1478

279 rows × 3 columns

About the datasets

```
In [16]: 1 print("Total Number of Images", flower_class.shape[0])
2 print("Number of Training Images", train1df.shape[0])
3 print("Number of Validation Images", val1df.shape[0])
4 print("Number of Training Images", test1df.shape[0])
```

```
Total Number of Images 1479
Number of Training Images 921
Number of Validation Images 279
Number of Training Images 279
```


Data Augmentation


```
In [17]: 1 datagen=ImageDataGenerator(rescale=1./255.,
2                               featurewise_center=False,
3                               featurewise_std_normalization=False,
4                               rotation_range=90,
5                               width_shift_range=0.2,
6                               height_shift_range=0.2,
7                               horizontal_flip=True,
8                               vertical_flip=True)
9
10 train1_generator=datagen.flow_from_dataframe(
11     dataframe=train1df,
12     directory="Flower/jpg/",
13     x_col="Images",
14     y_col="Class",
15     subset="training",
16     batch_size=128,
17     seed=42,
18     shuffle=True,
19     class_mode="sparse",
20     color_mode="rgb",
21     #target_size=(256,256))
22     target_size=(128,128))
23     print('Train generator created')
24
25
26 val1_generator=datagen.flow_from_dataframe(
27     dataframe=val1df,
28     directory="Flower/jpg/",
29     x_col="Images",
30     y_col="Class",
31     subset="training",
32     batch_size=138,
33     seed=42,
34     shuffle=True,
35     class_mode="sparse",
36     #target_size=(256,256))
37     target_size=(128,128))
38     print('Validation generator created')
39
40 test1_generator=datagen.flow_from_dataframe(
41     dataframe=test1df,
42     directory="Flower/jpg/",
43     x_col="Images",
44     y_col="Class",
45     subset="training",
46     #batch_size=128,
47     batch_size= 1,
48     color_mode="rgb",
49     seed=42,
50     shuffle=False,
51     #class_mode="sparse",
52     class_mode = None,
53     #target_size=(256,256))
54     target_size=(128,128))
55     print('Test generator created')
```

Found 921 validated image filenames belonging to 30 classes.

```
Train generator created
Found 279 validated image filenames belonging to 30 classes.
Validation generator created
Found 279 validated image filenames.
Test generator created
```

```
In [18]: 1 image, label = next(iter(train1_generator))
        2 plt.imshow(image[0,:]);
```



```
In [19]: 1 # plt.figure(figsize=(10,10))
        2 # for i in range(10):
        3 #     plt.subplot(5,5,i+1)
        4 #     plt.xticks([])
        5 #     plt.yticks([])
        6 #     plt.grid(False)
        7 #     plt.imshow(train_generator[i])
        8 train1_generator.image_shape
```

```
Out[19]: (128, 128, 3)
```

Create And Train the Model

```
In [20]: 1 model = Sequential()
2 #model.add(Conv2D(32, (3, 3), padding='same', input_shape=(256,256,3)))
3 model.add(Conv2D(32, (3, 3), padding='same', input_shape=(128,128,3)))
4 #model.add(BatchNormalization())
5 model.add(Activation('relu'))
6
7 model.add(Conv2D(32, (3, 3)))
8 #model.add(BatchNormalization())
9 model.add(Activation('relu'))
10 model.add(MaxPooling2D(pool_size=(2, 2)))
11 model.add(Dropout(0.25))
12
13 model.add(Conv2D(64, (3, 3), padding='same'))
14 #model.add(BatchNormalization())
15 model.add(Activation('relu'))
16
17 model.add(Conv2D(64, (3, 3)))
18 #model.add(BatchNormalization())
19 model.add(Activation('relu'))
20 model.add(MaxPooling2D(pool_size=(2, 2)))
21 model.add(Dropout(0.25))
22
23 model.add(Flatten())
24 model.add(Dense(512))
25 #model.add(BatchNormalization())
26 model.add(Activation('relu'))
27 model.add(Dropout(0.5))
28 #output Classes are 30
29 model.add(Dense(30))
30 model.add(Activation('softmax'))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

In [21]: 1 model.summary()

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 128, 128, 32)	896
activation_1 (Activation)	(None, 128, 128, 32)	0
conv2d_2 (Conv2D)	(None, 126, 126, 32)	9248
activation_2 (Activation)	(None, 126, 126, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 63, 63, 32)	0
dropout_1 (Dropout)	(None, 63, 63, 32)	0
conv2d_3 (Conv2D)	(None, 63, 63, 64)	18496
activation_3 (Activation)	(None, 63, 63, 64)	0
conv2d_4 (Conv2D)	(None, 61, 61, 64)	36928
activation_4 (Activation)	(None, 61, 61, 64)	0
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 64)	0
dropout_2 (Dropout)	(None, 30, 30, 64)	0
flatten_1 (Flatten)	(None, 57600)	0
dense_1 (Dense)	(None, 512)	29491712
activation_5 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 30)	15390
activation_6 (Activation)	(None, 30)	0
=====		
Total params: 29,572,670		
Trainable params: 29,572,670		
Non-trainable params: 0		
=====		

```
In [22]: 1 model.compile(optimizer = tf.train.AdamOptimizer(),
2               loss='sparse_categorical_crossentropy',
3               metrics=['accuracy'])
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3622: The name tf.log is deprecated. Please use tf.math.log instead.

Fitting the Model

```
In [33]: 1 from keras.callbacks import EarlyStopping
2
3 early_stopping_monitor = EarlyStopping(patience=3)
4 EPOCHS = 500
5
6 # STEP_SIZE_TRAIN=train1_generator.n//train1_generator.batch_size
7 # STEP_SIZE_VALID=val1_generator.n//val1_generator.batch_size
8 # STEP_SIZE_TEST=test1_generator.n//test1_generator.batch_size
9
10 history = model.fit(train1_generator,
11                    validation_data = val1_generator,
12                    epochs=EPOCHS, callbacks=[early_stopping_monitor])
```

Epoch 1/500

8/8 [=====] - 19s 2s/step - loss: 0.8202 - acc: 0.7088 - val_loss: 0.8842 - val_acc: 0.7168

Epoch 2/500

8/8 [=====] - 19s 2s/step - loss: 0.7804 - acc: 0.7161 - val_loss: 0.8556 - val_acc: 0.7276

Epoch 3/500

8/8 [=====] - 19s 2s/step - loss: 0.7301 - acc: 0.7567 - val_loss: 0.8404 - val_acc: 0.7204

Epoch 4/500

8/8 [=====] - 19s 2s/step - loss: 0.7286 - acc: 0.7446 - val_loss: 0.8485 - val_acc: 0.7025

Epoch 5/500

8/8 [=====] - 19s 2s/step - loss: 0.7095 - acc: 0.7357 - val_loss: 1.1007 - val_acc: 0.6380



Epoch 6/500

8/8 [=====] - 19s 2s/step - loss: 0.7913 - acc: 0.7248 - val_loss: 0.8823 - val_acc: 0.7133

```

In [39]: 1 def plot_history(history):
2     plt.figure(figsize=(15, 5))
3     plt.subplot(1, 2, 1)
4     plt.xlabel('Epochs')
5     plt.ylabel('Accuracy')
6     plt.plot(history.epoch, np.array(history.history['acc']),
7             label='Train')
8     plt.plot(history.epoch, np.array(history.history['val_acc']),
9             label='Val')
10    plt.title('Model Accuracy')
11    plt.legend()
12    plt.ylim([0, 2])
13
14    plt.subplot(1, 2, 2)
15    plt.plot(history.history['loss'], label='Train')
16    plt.plot(history.history['val_loss'], label='Val')
17    plt.title('Model Loss')
18    plt.ylabel('Loss')
19    plt.xlabel('Epochs')
20    plt.legend()
21    plt.show()
22
23    plot_history(history)
24
25

```



```

In [40]: 1 model.evaluate_generator(generator=val_generator)

```

```

Out[40]: [0.8677034198596913, 0.7096774280071259]

```

Predict Output

```
In [0]: 1 # Get the filenames from the generator
2 fnames = test1_generator.filenames
3
4 # Get the ground truth from generator
5 ground_truth = vall_generator.classes
6
7 # Get the label to class mapping from the generator
8 label2index = vall_generator.class_indices
9
10 # Getting the mapping from class index to class label
11 idx2label = dict((v,k) for k,v in label2index.items())
12
```

```
In [0]: 1 # Get the predictions from the model using the generator
2 predictions = model.predict_generator(vall_generator,verbose=1)
3 predicted_classes = np.argmax(predictions,axis=1)
```

```
In [67]: 1 errors = np.where(predicted_classes != ground_truth)[0]
2 print("No of errors = {}/{}".format(len(errors),vall_generator.samples))
```

No of errors = 265/279

```
In [81]: 1 # Show the errors
2 for i in range(len(errors)):
3     pred_class = np.argmax(predictions[errors[i]])
4     pred_label = idx2label[pred_class]
5
6     title = 'Original label:{}, Prediction :{}, confidence : {:.3f}'.format(
7         fnames[errors[i]].split('/')[0],
8         pred_label,
9         predictions[errors[i]][pred_class])
10
11     original = load_img('{}{}'.format(vall_generator.directory, fnames[errors[i]]))
12     plt.figure(figsize=[7,7])
13     plt.axis('off')
14     plt.title(title)
15     plt.imshow(original)
16     plt.show()
17
```

Original label:JFT_00002.jpg, Prediction :2, confidence : 0.977

